

Spatial factor analysis

James T. Thorson

```
library(tinyVAST)
library(fmesher)
set.seed(101)
```

tinyVAST is an R package for fitting vector autoregressive spatio-temporal (VAST) models. We here explore the capacity to specify a spatial factor analysis, where the spatial pattern for multiple variables is described via their estimated association with a small number of spatial latent variables.

Spatial factor analysis

We first explore the ability to specify two latent variables for five manifest variables. To start we simulate two spatial latent variables, project via a simulated loadings matrix, and then simulate a Tweedie response for each manifest variable:

```
# Simulate settings
theta_xy = 0.4
n_x = n_y = 10
n_c = 5
rho = 0.8
resid_sd = 0.5

# Simulate GMRFs
R_s = exp(-theta_xy * abs(outer(1:n_x, 1:n_y, FUN="-"))) )
R_ss = kronecker(X=R_s, Y=R_s)
delta_fs = mvtnorm::rmvnorm(n_c, sigma=R_ss )

#
L_cf = matrix( rnorm(n_c^2), nrow=n_c )
L_cf[,3:5] = 0
L_cf = L_cf + resid_sd * diag(n_c)

#
d_cs = L_cf %*% delta_fs
```

Where we can inspect the simulated loadings matrix

```
dimnames(L_cf) = list( paste0("Var ", 1:nrow(L_cf)),
                       paste0("Factor ", 1:ncol(L_cf)) )
knitr::kable( L_cf,
              digits=2, caption="True loadings")
```

Table 1: True loadings

	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5
Var 1	-0.77	0.26	0.0	0.0	0.0
Var 2	-0.66	1.03	0.0	0.0	0.0
Var 3	-0.30	-0.54	0.5	0.0	0.0
Var 4	-0.57	0.34	0.0	0.5	0.0
Var 5	-0.03	-0.24	0.0	0.0	0.5

We then specify the model as expected by *tinyVAST*:

```
# Shape into longform data-frame and add error
Data = data.frame( expand.grid(species=1:n_c, x=1:n_x, y=1:n_y),
                    "var"="logn", "z"=exp(as.vector(d_cs)) )
Data$n = tweedie::rtweedie( n=nrow(Data), mu=Data$z, phi=0.5, power=1.5 )
mean(Data$n==0)
#> [1] 0.03

# make mesh
mesh = fm_mesh_2d( Data[,c('x','y')] )

#
sem = "
  f1 -> 1, 11
  f1 -> 2, 12
  f1 -> 3, 13
  f1 -> 4, 14
  f1 -> 5, 15
  f2 -> 2, 16
  f2 -> 3, 17
  f2 -> 4, 18
  f2 -> 5, 19
  f1 <-> f1, NA, 1
  f2 <-> f2, NA, 1
  1 <-> 1, NA, 0
  2 <-> 2, NA, 0
  3 <-> 3, NA, 0
  4 <-> 4, NA, 0
  5 <-> 5, NA, 0
"

# fit model
out = fit( sem = sem,
           data = Data,
           formula = n ~ 0 + factor(species),
           spatial_graph = mesh,
           family = tweedie(),
           variables = c( "f1", "f2", 1:n_c ),
           data_colnames = list( space = c("x","y"), variable = "species",
                                time = "time", distribution = "dist"),
           control = tinyVASTcontrol(quiet=TRUE, trace=0, gmrf="proj") )

out
#> $call
```

```

#> fit(data = Data, formula = n ~ 0 + factor(species), sem = sem,
#>       family = tweedie(), data_colnames = list(space = c("x", "y"),
#>         variable = "species", time = "time", distribution = "dist"),
#>       variables = c("f1", "f2", 1:n_c), spatial_graph = mesh, control = tinyVASTcontrol(quiet = TRUE,
#>         trace = 0, gmrf = "proj"))
#>
#> $opt
#> $opt$par
#>      alpha_j      alpha_j      alpha_j      alpha_j      alpha_j      theta_z      theta_z      theta_z
#> 0.07570783 -0.02014888 0.22318277 0.14728087 -0.26514652 0.68016356 0.68285927 0.31701846 0.5
#>      log_sigma      log_sigma      log_kappa
#> -0.52205331 0.21851154 -0.26761196
#>
#> $opt$objective
#> [1] 631.3721
#>
#> $opt$convergence
#> [1] 0
#>
#> $opt$iterations
#> [1] 71
#>
#> $opt$evaluations
#> function gradient
#>      84      71
#>
#> $opt$message
#> [1] "relative convergence (4)"
#>
#>
#> $sdrep
#> sdreport(.) result
#>           Estimate Std. Error
#> alpha_j      0.07570783 0.31850774
#> alpha_j     -0.02014888 0.39763412
#> alpha_j      0.22318277 0.21847161
#> alpha_j      0.14728087 0.27057852
#> alpha_j     -0.26514652 0.14638317
#> theta_z      0.68016356 0.11510781
#> theta_z      0.68285927 0.15773444
#> theta_z      0.31701846 0.10358197
#> theta_z      0.52123769 0.10914344
#> theta_z      0.14819781 0.09200614
#> theta_z      0.51873790 0.13709521
#> theta_z     -0.31998633 0.10049164
#> theta_z      0.23601680 0.10640963
#> theta_z     -0.21613586 0.09652980
#> log_sigma -0.52205331 0.06761637
#> log_sigma  0.21851154 0.13313019
#> log_kappa -0.26761196 0.21030826
#> Maximum gradient component: 0.002233932
#>
#> $run_time

```

```
#> Time difference of 2.913271 secs
```

We can compare the true loadings (rotated to optimize comparison):

```
Lrot_cf = rotate_pca( L_cf )$L_tf
dimnames(Lrot_cf) = list( paste0("Var ", 1:nrow(Lrot_cf)),
                           paste0("Factor ", 1:ncol(Lrot_cf)) )
knitr::kable( Lrot_cf,
               digits=2, caption="Rotated true loadings")
```

Table 2: Rotated true loadings

	Factor 1	Factor 2	Factor 3	Factor 4	Factor 5
Var 1	-0.71	0.34	0.00	-0.11	-0.19
Var 2	-1.20	-0.18	0.00	-0.18	0.12
Var 3	0.22	0.73	-0.17	-0.09	0.10
Var 4	-0.70	0.25	0.06	0.37	0.03
Var 5	0.17	0.23	0.47	-0.08	0.03

with the estimated loadings

```
# Extract and rotate estimated loadings
Lhat_cf = matrix( 0, nrow=n_c, ncol=2 )
Lhat_cf[lower.tri(Lhat_cf,diag=TRUE)] = as.list(out$sdrop, what="Estimate")$theta_z
Lhat_cf = rotate_pca( L_tf=Lhat_cf, order="decreasing" )$L_tf
#> Warning in sqrt(Eigen$values): NaNs produced
```

Where we can compared the estimated and true loadings matrices:

```
dimnames(Lhat_cf) = list( paste0("Var ", 1:nrow(Lhat_cf)),
                           paste0("Factor ", 1:ncol(Lhat_cf)) )
knitr::kable( Lhat_cf,
               digits=2, caption="Rotated estimated loadings" )
```

Table 3: Rotated estimated loadings

	Factor 1	Factor 2
Var 1	0.64	-0.23
Var 2	0.82	0.26
Var 3	0.19	-0.41
Var 4	0.57	0.05
Var 5	0.07	-0.25

Or we can specify the model while ensuring that residual spatial variation is also captured:

```
#
sem = "
  f1 -> 1, 11
```

```

f1 -> 2, 12
f1 -> 3, 13
f1 -> 4, 14
f1 -> 5, 15
f2 -> 2, 16
f2 -> 3, 17
f2 -> 4, 18
f2 -> 5, 19
f1 <-> f1, NA, 1
f2 <-> f2, NA, 1
1 <-> 1, sd_resid
2 <-> 2, sd_resid
3 <-> 3, sd_resid
4 <-> 4, sd_resid
5 <-> 5, sd_resid
"

# fit model
out = fit( sem = sem,
           data = Data,
           formula = n ~ 0 + factor(species),
           spatial_graph = mesh,
           family = list( "obs"=tweedie() ),
           variables = c( "f1", "f2", 1:n_c ),
           data_colnames = list( space = c("x","y"), variable = "species",
                                time = "time", distribution = "dist"),
           control = tinyVASTcontrol(quiet=TRUE, trace=0, gmrf="proj") )

# Extract and rotate estimated loadings
Lhat_cf = matrix( 0, nrow=n_c, ncol=2 )
Lhat_cf[lower.tri(Lhat_cf,diag=TRUE)] = as.list(out$sdrep, what="Estimate")$theta_z
#> Warning in Lhat_cf[lower.tri(Lhat_cf, diag = TRUE)] = as.list(out$sdrep, : number of items to replace
Lhat_cf = rotate_pca( L_tf=Lhat_cf, order="decreasing" )$L_tf
#> Warning in sqrt(Eigen$values): NaNs produced

```

Where we can again compared the estimated and true loadings matrices:

```

dimnames(Lhat_cf) = list( paste0("Var ", 1:nrow(Lhat_cf)),
                          paste0("Factor ", 1:ncol(Lhat_cf)) )
knitr::kable( Lhat_cf,
              digits=2, caption="Rotated estimated loadings with full rank" )

```

Table 4: Rotated estimated loadings with full rank

	Factor 1	Factor 2
Var 1	0.69	-0.18
Var 2	0.74	0.23
Var 3	0.07	-0.42
Var 4	0.47	-0.02
Var 5	0.07	-0.07